

Minimum Spanning Trees



© 2013 Goodrich, Tamassia, Goldwasser

Minimum Spanning Trees

1

Minimum Spanning Trees

Spanning subgraph

- Subgraph of a graph G containing all the vertices of G

Spanning tree

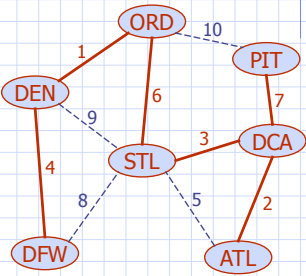
- Spanning subgraph that is itself a (free) tree

Minimum spanning tree (MST)

- Spanning tree of a weighted graph with minimum total edge weight

Applications

- Communications networks
- Transportation networks



© 2013 Goodrich, Tamassia, Goldwasser

Minimum Spanning Trees

2

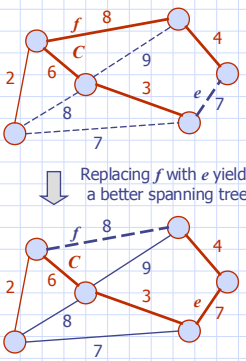
Cycle Property

Cycle Property:

- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T and C let be the cycle formed by e with T
- For every edge f of C , $weight(f) \leq weight(e)$

Proof:

- By contradiction
- If $weight(f) > weight(e)$ we can get a spanning tree of smaller weight by replacing e with f



© 2013 Goodrich, Tamassia, Goldwasser

Minimum Spanning Trees

3

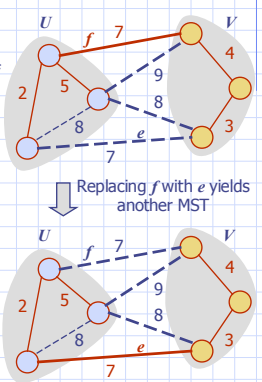
Partition Property

Partition Property:

- Consider a partition of the vertices of G into subsets U and V
- Let e be an edge of minimum weight across the partition
- There is a minimum spanning tree of G containing edge e

Proof:

- Let T be an MST of G
- If T does not contain e , consider the cycle C formed by e with T and let f be an edge of C across the partition
- By the cycle property, $weight(f) \leq weight(e)$
- Thus, $weight(f) = weight(e)$
- We obtain another MST by replacing f with e



© 2013 Goodrich, Tamassia, Goldwasser

Minimum Spanning Trees

4


Prim-Jarnik's Algorithm

- Similar to Dijkstra's algorithm
- We pick an arbitrary vertex s and we grow the MST as a cloud of vertices, starting from s
- We store with each vertex v label $d(v)$ representing the smallest weight of an edge connecting v to a vertex in the cloud
- At each step:
 - We add to the cloud the vertex u outside the cloud with the smallest distance label
 - We update the labels of the vertices adjacent to u


Prim-Jarnik Pseudo-code

```
Algorithm PrimJarnik( $G$ ):
  Input: An undirected, weighted, connected graph  $G$  with  $n$  vertices and  $m$  edges
  Output: A minimum spanning tree  $T$  for  $G$ 
  Pick any vertex  $s$  of  $G$ 
   $D[s] = 0$ 
  for each vertex  $v \neq s$  do
     $D[v] = \infty$ 
  Initialize  $T = \emptyset$ .
  Initialize a priority queue  $Q$  with an entry  $(D[v], (v, \text{None}))$  for each vertex  $v$ ,
  where  $D[v]$  is the key in the priority queue, and  $(v, \text{None})$  is the associated value.
  while  $Q$  is not empty do
     $(u, e) = \text{value returned by } Q.\text{remove.min}()$ 
    Connect vertex  $u$  to  $T$  using edge  $e$ .
    for each edge  $e' = (u, v)$  such that  $v$  is in  $Q$  do
      {check if edge  $(u, v)$  better connects  $v$  to  $T$ }
      if  $w(u, v) < D[v]$  then
         $D[v] = w(u, v)$ 
        Change the key of vertex  $v$  in  $Q$  to  $D[v]$ .
        Change the value of vertex  $v$  in  $Q$  to  $(v, e')$ .
  return the tree  $T$ 
```

Example



Example (contd.)



Analysis

- Graph operations
 - We cycle through the incident edges once for each vertex
- Label operations
 - We set/get the distance, parent and locator labels of vertex z : $O(\deg(z))$ times
 - Setting/getting a label takes $O(1)$ time
- Priority queue operations
 - Each vertex is inserted once into and removed once from the priority queue, where each insertion or removal takes $O(\log n)$ time
 - The key of a vertex w in the priority queue is modified at most $\deg(w)$ times, where each key change takes $O(\log n)$ time
- Prim-Jarnik's algorithm runs in $O((n + m) \log n)$ time provided the graph is represented by the adjacency list structure
 - Recall that $\sum_v \deg(v) = 2m$
- The running time is $O(m \log n)$ since the graph is connected

Python Implementation

```

1 def MST_PrimJarnik(g):
2     """Compute a minimum spanning tree of weighted graph g.
3     """
4     Return a list of edges that comprise the MST (in arbitrary order).
5
6     d = {} # d[v] is bound on distance to tree
7     tree = [] # list of edges in spanning tree
8     pq = AdaptableHeapPriorityQueue() # d[v] maps to value (v, e=(u,v))
9     plocator = {} # map from vertex to its pq locator
10
11 # for each vertex v of the graph, add an entry to the priority queue, with
12 # the source having distance 0 and all others having infinite distance
13 for v in g.vertices():
14     if len(d) == 0: # this is the first node
15         d[v] = 0 # make it the root
16     else:
17         d[v] = float('inf') # positive infinity
18     plocator[v] = pq.add(d[v], (v, None))
19
20 while not pq.is_empty():
21     key, value = pq.remove_min()
22     u, edge = value # unpack tuple from pq
23     del plocator[u] # u is no longer in pq
24     if edge is not None: # add edge to tree
25         tree.append(edge)
26     for link in g.incident_edges(u):
27         v = link.opposite(u) # thus v not yet in tree
28         if v in plocator:
29             # see if edge (u,v) better connects v to the growing tree
30             wgt = link.element()
31             if wgt < d[v]: # better edge to v?
32                 d[v] = wgt # update the distance
33                 pq.update(plocator[v], d[v], (v, link)) # update the pq entry
34 return tree

```

Kruskal's Approach

- Maintain a partition of the vertices into clusters
 - Initially, single-vertex clusters
 - Keep an MST for each cluster
 - Merge "closest" clusters and their MSTs
- A priority queue stores the edges outside clusters
 - Key: weight
 - Element: edge
- At the end of the algorithm
 - One cluster and one MST

Kruskal's Algorithm

Algorithm Kruskal(G):

Input: A simple connected weighted graph G with n vertices and m edges

Output: A minimum spanning tree T for G

for each vertex v in G do

Define an elementary cluster $C(v) = \{v\}$.

Initialize a priority queue Q to contain all edges in G , using the weights as keys.

$T = \emptyset$ { T will ultimately contain the edges of the MST}

while T has fewer than $n - 1$ edges do

$(u, v) = \text{value returned by } Q.\text{remove_min}()$



Let $C(u)$ be the cluster containing u , and let $C(v)$ be the cluster containing v .

if $C(u) \neq C(v)$ then

Add edge (u, v) to T .

Merge $C(u)$ and $C(v)$ into one cluster.

return tree T



Data Structure for Kruskal's Algorithm

- The algorithm maintains a forest of trees
- A priority queue extracts the edges by increasing weight
- An edge is accepted if it connects distinct trees
- We need a data structure that maintains a **partition**, i.e., a collection of disjoint sets, with operations:
 - **makeSet**(u): create a set consisting of u
 - **find**(u): return the set storing u
 - **union**(A, B): replace sets A and B with their union

© 2013 Goodrich, Tamassia, Goldwasser Minimum Spanning Trees 15

List-based Partition

- Each set is stored in a sequence
- Each element has a reference back to the set
 - operation **find**(u) takes $O(1)$ time, and returns the set of which u is a member.
 - in operation **union**(A,B), we move the elements of the smaller set to the sequence of the larger set and update their references
 - the time for operation **union**(A,B) is $\min(|A|, |B|)$
- Whenever an element is processed, it goes into a set of size at least double, hence each element is processed at most $\log n$ times

© 2013 Goodrich, Tamassia, Goldwasser Minimum Spanning Trees 16

Partition-Based Implementation

- Partition-based version of Kruskal’s Algorithm
 - Cluster merges as unions
 - Cluster locations as finds
- Running time $O((n + m) \log n)$
 - Priority Queue operations: $O(m \log n)$
 - Union-Find operations: $O(n \log n)$

Python Implementation

```
1 def MST_Kruskal(g):
2     """ Compute a minimum spanning tree of a graph using Kruskal's algorithm.
3
4     Return a list of edges that comprise the MST.
5
6     The elements of the graph's edges are assumed to be weights.
7
8     """
9     tree = []
10    pq = HeapPriorityQueue()
11    forest = Partition()
12    position = {}
13
14    for v in g.vertices():
15        position[v] = forest.make_group(v)
16
17    for e in g.edges():
18        pq.add(e.element(), e)
19
20    size = g.vertex_count()
21    while len(tree) != size - 1 and not pq.is_empty():
22        weight, edge = pq.remove_min()
23        u, v = edge.endpoints()
24        a = forest.find(position[u])
25        b = forest.find(position[v])
26        if a != b:
27            tree.append(edge)
28            forest.union(a, b)
29
30    return tree
```

Baruvka’s Algorithm (Exercise)

- Like Kruskal’s Algorithm, Baruvka’s algorithm grows many clusters at once and maintains a forest T
- Each iteration of the while loop halves the number of connected components in forest T
- The running time is $O(m \log n)$

Algorithm *BaruvkaMST*(G)

$T \leftarrow V$ {just the vertices of G }

while T has fewer than $n - 1$ edges do

for each connected component C in T do

Let edge e be the smallest-weight edge from C to another component in T

if e is not already in T then

Add edge e to T

return T

Example of Baruvka’s Algorithm (animated)

Slide by Matt Stallmann included with permission.

