

Shortest Paths

© 2013 Goodrich, Tamassia, Goldwasser

Shortest Paths

1

Weighted Graphs

- In a weighted graph, each edge has an associated numerical value, called the weight of the edge
- Edge weights may represent, distances, costs, etc.
- Example:
 - In a flight route graph, the weight of an edge represents the distance in miles between the endpoint airports

© 2013 Goodrich, Tamassia, Goldwasser

Shortest Paths

2

Shortest Paths

- Given a weighted graph and two vertices u and v , we want to find a path of minimum total weight between u and v .
 - Length of a path is the sum of the weights of its edges.
- Example:
 - Shortest path between Providence and Honolulu
- Applications
 - Internet packet routing
 - Flight reservations
 - Driving directions

© 2013 Goodrich, Tamassia, Goldwasser

Shortest Paths

3

Shortest Path Properties

- Property 1:**
 - A subpath of a shortest path is itself a shortest path
- Property 2:**
 - There is a tree of shortest paths from a start vertex to all the other vertices

Example:
Tree of shortest paths from Providence

© 2013 Goodrich, Tamassia, Goldwasser

Shortest Paths


4

Dijkstra's Algorithm


- The distance of a vertex v from a vertex s is the length of a shortest path between s and v
- Dijkstra's algorithm computes the distances of all the vertices from a given start vertex s
- Assumptions:
 - the graph is connected
 - the edges are undirected
 - the edge weights are **nonnegative**
- We grow a "cloud" of vertices, beginning with s and eventually covering all the vertices
- We store with each vertex v a label $d(v)$ representing the distance of v from s in the subgraph consisting of the cloud and its adjacent vertices
- At each step
 - We add to the cloud the vertex u outside the cloud with the smallest distance label, $d(u)$
 - We update the labels of the vertices adjacent to u

Edge Relaxation


- Consider an edge $e = (u, z)$ such that
 - u is the vertex most recently added to the cloud
 - z is not in the cloud
- The relaxation of edge e updates distance $d(z)$ as follows:
$$d(z) \leftarrow \min\{d(z), d(u) + \text{weight}(e)\}$$



Example



Example (cont.)



Dijkstra's Algorithm

Algorithm ShortestPath(G, s):
Input: A weighted graph G with nonnegative edge weights, and a distinguished vertex s of G .
Output: The length of a shortest path from s to v for each vertex v of G .
Initialize $D[s] = 0$ and $D[v] = \infty$ for each vertex $v \neq s$.
Let a priority queue Q contain all the vertices of G using the D labels as keys.
while Q is not empty **do**
 {pull a new vertex u into the cloud}
 $u =$ value returned by $Q.remove_min()$
 for each vertex v adjacent to u such that v is in Q **do**
 {perform the *relaxation* procedure on edge (u, v) }
 if $D[u] + w(u, v) < D[v]$ **then**
 $D[v] = D[u] + w(u, v)$
 Change to $D[v]$ the key of vertex v in Q .
return the label $D[v]$ of each vertex v

Analysis of Dijkstra's Algorithm


- Graph operations
 - We find all the incident edges once for each vertex
- Label operations
 - We set/get the distance and locator labels of vertex z $O(\deg(z))$ times
 - Setting/getting a label takes $O(1)$ time
- Priority queue operations
 - Each vertex is inserted once into and removed once from the priority queue, where each insertion or removal takes $O(\log n)$ time
 - The key of a vertex in the priority queue is modified at most $\deg(u)$ times, where each key change takes $O(\log n)$ time
- Dijkstra's algorithm runs in $O((n + m) \log n)$ time provided the graph is represented by the adjacency list/map structure
 - Recall that $\sum_v \deg(v) = 2m$
- The running time can also be expressed as $O(m \log n)$ since the graph is connected

Python Implementation

```
1 def shortest_path_lengths(g, src):
2     """Compute shortest-path distances from src to reachable vertices of g.
3
4     Graph g can be undirected or directed, but must be weighted such that
5     e.element() returns a numeric weight for each edge e.
6
7     Returns dictionary mapping each reachable vertex to its distance from src.
8     """
9     d = {} # d[v] is upper bound from s to v
10    cloud = {} # map reachable v to its d[v] value
11    pq = AdjHeapPriorityQueue() # vertex v will have key d[v]
12    pe_locator = {} # map from vertex to its pq locator
13
14    # for each vertex v of the graph, add an entry to the priority queue, with
15    # the source having distance 0 and all others having infinite distance
16    for v in g.vertices():
17        if v is src:
18            d[v] = 0
19        else:
20            d[v] = float('inf') # syntax for positive infinity
21            pe_locator[v] = pq.add(d[v], v) # save locator for future updates
22
23    while not pq.is_empty():
24        key, u = pq.remove_min()
25        cloud[u] = key # its correct d[u] value
26        del pe_locator[u] # u is no longer in pq
27        for e in g.incident_edges(u): # outgoing edges (u,v)
28            v = e.opposite(u)
29            if v not in cloud:
30                # perform relaxation step on edge (u,v)
31                wgt = e.element()
32                if d[u] + wgt < d[v]: # better path to v?
33                    d[v] = d[u] + wgt # update the distance
34                    pq.update(pe_locator[v], d[v], v) # update the pq entry
35
36    return cloud # only includes reachable vertices
```


Why Dijkstra's Algorithm Works

- Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.
 - Suppose it didn't find all shortest distances. Let F be the first wrong vertex the algorithm processed.
 - When the previous node, D , on the true shortest path was considered, its distance was correct
 - But the edge (D, F) was *relaxed* at that time!
 - Thus, so long as $d(F) \geq d(D)$, F 's distance cannot be wrong. That is, there is no wrong vertex



Why It Doesn't Work for Negative-Weight Edges

- ◆ Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.
- If a node with a negative incident edge were to be added late to the cloud, it could mess up distances for vertices already in the cloud.




Bellman-Ford Algorithm (not in book)

- Works even with negative-weight edges
- Must assume directed edges (for otherwise we would have negative-weight cycles)
- Iteration i finds all shortest paths that use i edges.
- Running time: $O(nm)$.
- Can be extended to detect a negative-weight cycle if it exists
 - How?

```
Algorithm BellmanFord( $G, s$ )
for all  $v \in G.vertices()$ 
  if  $v = s$ 
     $setDistance(v, 0)$ 
  else
     $setDistance(v, \infty)$ 
for  $i \leftarrow 1$  to  $n - 1$  do
  for each  $e \in G.edges()$ 
    { relax edge  $e$  }
     $u \leftarrow G.origin(e)$ 
     $z \leftarrow G.opposite(u, e)$ 
     $r \leftarrow getDistance(u) + weight(e)$ 
    if  $r < getDistance(z)$ 
       $setDistance(z, r)$ 
```

Bellman-Ford Example

Nodes are labeled with their $d(v)$ values



DAG-based Algorithm (not in book)

- Works even with negative-weight edges
- Uses topological order
- Doesn't use any fancy data structures
- Is much faster than Dijkstra's algorithm
- Running time: $O(n+m)$.

```
Algorithm DagDistances( $G, s$ )
for all  $v \in G.vertices()$ 
  if  $v = s$ 
     $setDistance(v, 0)$ 
  else
     $setDistance(v, \infty)$ 
{ Perform a topological sort of the vertices }
for  $u \leftarrow 1$  to  $n$  do {in topological order}
  for each  $e \in G.outEdges(u)$ 
    { relax edge  $e$  }
     $z \leftarrow G.opposite(u, e)$ 
     $r \leftarrow getDistance(u) + weight(e)$ 
    if  $r < getDistance(z)$ 
       $setDistance(z, r)$ 
```

